Copyright Model for Collaboration
Literature Review

William Marsey
Imperial College London

June 2014

1 Introduction

1.1 The Research Question

Given a collaboratively edited document, and information about the collection of edits that constitute that document, may we measure the quality of each contribution? And may we use that to give all the collaborators a algorithmically-defined 'stake' in that final document?

Collaborative work is becoming a big deal. It is both interesting and an important trend in modern computer use. And the data is abundant.

Amongst many other things, this topic is a playground for sociology, machine learning, network studies, as well as more general studies of conflict, and personality. My work intends to focus on the algorithmic side of things - approximate string matching in particular. I look at how we may use Levenshtein distance, and the various favours, varieties and optimizations thereof, to measure contribution to a collaborative text, and how we may implement a version of this algorithm specifically tailored to our needs.

The main questions we ask are:

- What does Levenshtein distance define of a contribution in the context of massive online collaboration?
- What are the limitations and implications of defining contribution in this way?
- What else may we learn from analysing contribution?

We base our studies around data from Wikipedia. This study is defined by – and in some ways determined by – the specific context of Wikipedia, but, as we will see, is ultimately enriched by it. Due to its open-source nature, and its size, studies that touch upon Wikipedia cover a very broad range of topics. Many of them are directly related to the topic we concern ourselves with here, and many more may enrich our study tangentially.

2 Previous work

There are three sections here for the three different topics that come to bear on this subject:

- Wikipedia, studies of wikipedia, and the nature of Wikipedia
- The 'edit distance problem', Levenshtein distance
- The various pre-existing studies that apply the latter to the former

2.1 On Wikipedia

Wikipedia's pre-eminence as an online resource is self-evident to anyone who has searched the internet for a generic topic. The website is ranked 6^{th} globally in terms of website traffic,[1] and is the highest-ranked reference website by far - most of the sites it shares the top spots with are portals, search engines, shopping mega-sites, and social media websites.[2] Despite some skepticism (particularly concern over the inherent chaos of the system: "...edits, contributed in a predominantly undirected and haphazard fashion by ... unvetted volunteers."[3]), it is widely claimed to be a success, 'the best-developed attempt thus far of the enduring quest to gather all human knowledge in one place'[4].

That Wikipedia has become a hub of research in many fields is also self-evident to anyone who has searched for articles on the subject. Mesgari et al, just quoted, has prepared a very recent 'systematic review of scholarly research on the content of Wikipedia', which gives an overview of 110 articles on the subject — an attestment to the observation that Wikipedia has been 'irresistable point of unquiry for researchers from various fields of knowledge', and a useful touching stone for this study. Mesgari et al's review finds 82 out of the 100 to concern quality in Wikipedia articles, some of these are also referenced here, and many of the others will come to bear on the study as it progresses.

Other important general sources will be WikiLit,[5] AcaWiki[6] and WikiPapers[7], all of which are online repositories of academic research into Wikipedia and other Wikis (as well as being Wikis themselves...).

2.1.1 The particularities of Wikipedia, and implications

The six 'risks' one takes when referencing Wikipedia, as defined in an early article on the subject,[8] is a good starting block for identifying the ways in which to ragard the 'quality' of content in Wikipedia. We list them here, describing the implications for our work with each. Some are particular to Wikipedia, some are inherent to all Wikis.

- **Accuracy.** It is important to remember that, without severely increasing the complexity of our algorithm, we may not verify the accuracy of information. And, if accuracy of information is proportionate to value (surely it must be in a reference text), then our algorithm may misplace its reward. We may most usefully look at the problem as follows. The texts that are edited most often are those that are visited most often. The previously cited studies of Lih and Mestyan et al attest to this - they both study the peaks in activity in articles that are brought to attention in some way. We find in the work of Bongwon Suh et al that the growth of wikipedia is inverse-exponential, as the overheads of coordination and beaurocrosy temper content creation.[9][10] Content is more likely to be refined and corrected as an article grows old.[11] We can assume, then, that all articles tend towards accuracy (we may find this bore out in Giles's 2005 semi-formal comparison of Encyclopedia Britannica articles to Wikipedia

---

[1] According to Alexa, an Amazon-owned company. The statistics are wide-rangingbased on a combined measure of Unique Visitors and Pageviews, and the data mined from around 25,000 different browser extensions, as well as sites that have installed Alexa's scripts. (*Alexa About Us*) Alexa may well be biased towards English speakers and Internet Explorer users, but this may underestimate Wikipedia.org's popularity, since 'two thirds of all Wikipedia articles are in languages other than English' (*Wikimedia Wikipedia.org is more popular than...: Note on Alexa rankings*)

[2] *Alexa The top 500 sites on the web.*

[3] Wilkinson and Huberman 2007.

[4] Mesgari et al. 2014.

[5] *WikiLit.*

[6] *AcaWiki.*

[7] *WikiPapers.*

[8] Denning et al. 2005.

[9] Suh et al. 2009.

[10] Kittur et al. 2007.

[11] Wilkinson and Huberman 2007.

articles - finding an average three mistakes in the former and 4 mistakes in the latter)[12]. We may possibly extend this to say that all edits improve a text. This is complicated by malicious, misinformed or malformed edits, but we will discuss dealing with these later.
make

- **Motives.** It has been found that different contributors may edit Wikipedia in various different ways, according to their proclivities.
  Response:
- **Uncertain Expertise.** We may take this to mean malformed and misinformed editing, but we may also take it to mean malicious editing. As for malicious edits - we find a lot of useful information in Potthast et al's work on automatic detection of vandals,[13] as well as the discussions around Wikipedia's own Counter-Vandalism Unit ('CVU').[14] including
- **Volatility.**
- **Coverage.** Cite that structure is a problem.
  Response: We may want to reward extra for restructuring.
- **Sources.** There has a been some work explicitly taking external links to be relative to quality,[15] and seems to have been bore out by a further study which found this to be a heuristic used by actual readers.[16].
  Response: We may give give extra weight to the value of (working) hyperlinks, and fixing hyperlinks.

2.2 Wikipedia Self-Valuation

> [Wikipedia] cannot attain the status of a true encyclopedia without more formal content-inclusion and expert review procedures.[17]

Most visited articles in an hour - correlates with (american-centric) events[18]

Denning says it cannot attain the status of a true encyclopedia without more formal content-inclusion and expert review procedures[19] this corroborates by findings in[20]?

2.2.1 On Wikipedia

'robust and remarkable growth'[21][22]

Wikipedia, at the last dump, consisted 800G of compressed data[23]

2.2.2 Evaluating Wikipedia articles

identify, analyse

after article mentioned in press[24]

---

[12]Giles 2005.

[13]Potthast, Stein, and Gerling 2008.

[14]*Wikipedia Counter-Vandalism Unit*.

[15]**CITEHYPERLINKS**.

[16]**THISHEURISTICIGUESS**.

[17]Denning et al. 2005.

[18]*Wikipedia Wikipedia Signpost/2013-02-04/Special report: Examining the popularity of Wikipedia articles: catalysts, trends, and applications*.

[19]Denning et al. 2005.

[20]Giles 2005.

[21]Kittur et al. 2007.

[22]Voss 2005.

[23]*Wikimedia enwiki dump progress on 20140502*.

[24]Lih 2004.

compared by 'experts' to 'equivalent' Encyclopedia Britannica articles[25]

found metrics of article quality through factor analysis[26]

Analysis by conflict - revisions?[27]

WikiTrust. The most 'complete' of the many of the. Exists as firefox plugin (though it doesn't work any more) Culmination of various studies that try to QUOTE[28] and QUOTE CITE. It was assessed as recently as 2011[29]

2.3 Edit difference

Standard: Levenshtein distance[30]

3 Previous Work

different views emerging topics gaps and inconsistencies

Goals.

Methods. Python. Levenshtein. Optimisations of.

For this study we will assume that the 'final' or 'target' version is fully trusted, or, that its quality does not need to be evaluated by us, and does not affect how we evaluate contributions. This way we can leave issues of trust regarding the article itself to one side, and concentrate on the various edits made. The reputation, or quality, of the article itself is not important for this study. Our intention is to evaluate an individual's weighted stake in an article - whether that article is of good or bad quality is something of a different issue.

However, some of the methods used to measure quality are directly applicable here, and, as mentioned previously, are copious. Of particular interest is the academic work that culminated in WikiTrust,[31][32] and other studies of significant subordinate importance.[33][34]

Wikitrust was[35] a firefox plugin, designed to highlight the words of a Wikipedia article with different gradations of yellow. The gradations relate to levels of trust, a screenshot can be seen below.

[36]

4 Conclusions

Possible extensions. Perhaps we can examine more about what we may find out about the articles. Cite other studies Lih's 2004 study of articles immediately after they have been cited in the press[37],

---

[25] Giles 2005.

[26] Stvilia et al. 2005.

[27] Kittur et al. 2007.

[28] Adler and Alfaro 2007.

[29] Lucassen and Schraagen 2011.

[30] Levenshtein 1966.

[31] De Alfaro et al. 2011.

[32] Adler and Alfaro 2007.

[33] Zeng et al. 2006.

[34] Cross 2006.

[35] Defunct as of author's checks, Apr 2014

[36] Lucassen and Schraagen 2011.

[37] Lih 2004.

Figure 1: Wikitrust in action (2011)

and Metyan's 2012 use of the site to predict box-office success[38]. Lieberman worked out their locale.[39]:

1. peaks in activity (rate of levenshtein distance... may have to be in a log graph...)
2.

PREDIFINED / NOT-PREDEFINED ideas of quality. look for when the the article levels off? And do this by DATE rather than REVISION. We may assume that pageviews are more well-distributed than revisions

summarize major contributions (which do we care about?) evaluate your current position point out any flaw in methodology/research/contradictions are there any gaps in the area which you will cover in your research? How will you integrate sources you have mentioned into your dissertation? Point out any areas for further study

$$ed(v_{i-1}, v_i) < ed(v_{i-1}, v_i) + ed(v_i, v_{i+1})$$

Figure 2: Diagram showing identification of a partially undone operation

We will also consider introducing ways of representing the history of a page in space. The position of a certain revision in that space may help us define the corresponding edit operation.

Since we are taking one revision, say, the current one, to be the ultimate 'target' of previous edits, perhaps we should describe all previous edits in terms of this final version. We may measure the edit distance between all previous versions and arrive at the set:

where v_n is the 'target' version of a page, and we are considering n revisions of the page

---

[38]Mestyán, Yasseri, and Kertász 2012.
[39]Lieberman and Lin 2009.

$$ed(v_{i-1}, v_{i+1}) = 0, \ ed(v_{i-1}, v_i) > 0)$$

Figure 3: Diagram showing identification of an undone operation

$$\text{prev} = < ed(v_0, v_n), ed(v_1, v_n), ..., ed(v_n, v_n) >$$

$$v_i$$

$$v_{i-1} \qquad\qquad\qquad v_{i+1}$$

$$ed(v_{i-1}, v_{i+1}) = 0, \ ed(v_{i-1}, v_i) > 0)$$

Figure 4: Diagram showing revisions

5 Progress report

The Code in appendix A

A   Appendix A: Code progress

A.1   Python class for scraping a Wikipedia article's revision history

The following code is a first draft of a class which incrementally traces, parses, and stores the revision history of select articles. It chooses random articles up to a maximum amount of articles unless an article (or articles) are specified. It traces the entire discoverable[40] history of the, unless a specific depth is specified by the user.

The class already yields workable data, but here is some immediate further work for this code:

- Allow the user to specify timeframe
- Allow for integration with a postgres database (at the moment the code saves the data in CSV format).

This code is an extension to an existing wikipedia API class for python (which did not provide the features we need here).[41]

```python
import requests
import time
import json
import csv
import wikipedia
from bs4 import BeautifulSoup
from datetime import datetime, timedelta
from decimal import Decimal

WIKI_API_URL = 'http://en.wikipedia.org/w/api.php'
WIKI_USER_AGENT = 'wikipedia (https://github.com/goldsmith/Wikipedia/)'

class WikiRevisionScrape:
    par = {
        'format': 'json',
        'action': 'query',
        'prop': 'revisions'
        }
    head = {
        'User-Agent': WIKI_USER_AGENT
        }
    rand = True
    pagelimit = 1
    historylimit = -1
    rl = False
    rl_minwait = None
    rl_lastcall = None
    pageid = 0
    parentid = 0
    childid = 0

    #atm naively assuming headers, params, titles to be in correct format
    def __init__(self, pagelimit=1, historylimit=-1, _headers=None, _params=None,
        _titles=None):
        if(_params):
            params = _params

        if(_headers):
```

---

[40]Using the Wikipedia API, articles can either be traced back to their origin (revision parent ID = 0), or to the point at which a loop is found in the revision history - this usually happens with older articles.

[41]*Python wikipedia 1.3.0*.

```python
        self.head = _headers

    if(_titles):
        self.params['titles'] = _titles
        self.rand = False

    self.pagelimit = pagelimit
    self.historylimit = historylimit


def scrape(self, indexfilename, contentsfilename):
    self._pace()
    index_f = open(indexfilename + ".csv", "ab") #HACK = needs to migrate to postrgres
    contents_f = open(contentsfilename + ".csv", "ab") #HACK = needs to migrate to
        postrgres
    index = csv.writer(index_f)
    contents = csv.writer(contents_f)
    index.writerow(["PAGEID","REVISION","USER","USERID","TIMSTAMP","SIZE","COMMENT"])
    contents.writerow(["PAGEID","REVISION","CONTENT"])

    for i in range(self.pagelimit):
        if 'rvprop' in self.par:
            del self.par['rvprop']
        if 'revids' in self.par:
            del self.par['revids']
        print "fetching page"
        if(self.rand):
            self.par['titles'] = wikipedia.random() #get random title
        self.childid = self._getlatest()
        r = requests.get(WIKI_API_URL, params=self.par, headers=self.head)
        self._rate()
        del self.par['titles']
        self._tracehist(index, contents)

def _getlatest(self):
    r = requests.get(WIKI_API_URL, params=self.par, headers=self.head)
    r = r.json()

    #HACK = should grab multiple pages
    for key, value in r['query']['pages'].iteritems():
        self.pageid = key
    #HACK = chould grab multiple revisions (for each pageid)
    self.parentid = self.childid =
        r['query']['pages'][self.pageid]['revisions'][0]['revid']
    return self.childid

def _tracehist(self, index, contents):
    ##We store revisions we've visited
    ##loops can occur in revision histories
    visited = []
    i = self.historylimit
    j = 0

    self.par['rvprop'] =
        'userid|user|ids|flags|tags|size|comment|contentmodel|timestamp|content'

    while (self.parentid not in visited) and i is not 0 and self.parentid is not 0:
        self.par['revids'] = self.parentid

        self._pace()
```

8

```python
        r = requests.get(WIKI_API_URL, params=self.par, headers=self.head)
        r = r.json()

        self._rate()

        visited.append(self.childid)

        #print r

        self.childid = r['query']['pages'][self.pageid]['revisions'][0]['revid']
        self.parentid = r['query']['pages'][self.pageid]['revisions'][0]['parentid']
        user = r['query']['pages'][self.pageid]['revisions'][0]['user']
        userid = r['query']['pages'][self.pageid]['revisions'][0]['userid']
        size = r['query']['pages'][self.pageid]['revisions'][0]['size']
        timestamp = r['query']['pages'][self.pageid]['revisions'][0]['timestamp']
        comment = "" #comments sometimes don't return from old revisions...
        try:
            comment = r['query']['pages'][self.pageid]['revisions'][0]['comment']
        except:
            comment = ""
        content = r['query']['pages'][self.pageid]['revisions'][0]['*']

        index.writerow([self.pageid, self.childid, user.encode("UTF-8"), userid,
            timestamp, size, comment.encode("UTF-8")])
        contents.writerow([self.pageid, self.childid, content.encode("UTF-8")])

        if(self.historylimit > 0):
            print self.pageid, "fetch", j+1, "of", self.historylimit, ", revid",
                self.childid, "timestamp", str(timestamp)
            i = i - 1
        else:
            print self.pageid, "fetch", j+1, ", revid", self.childid, "timestamp",
                str(timestamp)
        j = j + 1
    print "limit reached"

def _pace(self):
    if self.rl and self.rl_last_call and self.rl_lastcall + self.rl_minwait >
        datetime.now():
        wait_time = (self.rl_lastcall + self.rl_minwait) - datetime.now()
        time.sleep(int(wait_time.total_seconds()))

def _rate(self):
    if self.rl:
        self.rl_lastcall = datetime.now()
```

A.2 Example output

```
$ python
>>> import WikiRevisionScrape

>>> singlescraper = WikiRevisionScrape.WikiRevisionScrape()

>>> singlescraper.scrape("filename1","filename2")
fetching page
searching for
{u'action': u'query', u'list': u'random', u'rnlimit': 1,
```

```
      u'rnnamespace': 0, u'format': u'json'}
25455543 fetch 1 , revid 553292956 timestamp 2013-05-03T03:01:26Z
25455543 fetch 2 , revid 550043052 timestamp 2013-04-12T18:59:57Z
25455543 fetch 3 , revid 503496279 timestamp 2012-07-21T21:52:51Z
.
. [skipping some output]
.
25455543 fetch 23 , revid 331902859 timestamp 2009-12-15T23:25:23Z
25455543 fetch 24 , revid 331902368 timestamp 2009-12-15T23:22:50Z
25455543 fetch 25 , revid 331902181 timestamp 2009-12-15T23:21:47Z
limit reached

>>> multiscraper = WikiRevisionScrape.WikiRevisionScrape(pagelimit=1000)

>>> multiscraper.scrape("multifilename1","multifilename2")
fetching page
searching for
{u'action': u'query', u'list': u'random', u'rnlimit': 1,
      u'rnnamespace': 0, u'format': u'json'}
7096591 fetch 1 , revid 472732138 timestamp 2012-01-23T03:00:01Z
7096591 fetch 2 , revid 416290467 timestamp 2011-02-28T00:06:47Z
.
. [skipping some output]
.
7096591 fetch 8 , revid 89546539 timestamp 2006-11-22T23:31:09Z
7096591 fetch 9 , revid 77039186 timestamp 2006-09-21T20:00:55Z
limit reached
fetching page
searching for
{u'action': u'query', u'list': u'random', u'rnlimit': 1,
      u'rnnamespace': 0, u'format': u'json'}
24830105 fetch 1 , revid 547881527 timestamp 2013-03-30T21:34:39Z
24830105 fetch 2 , revid 500160388 timestamp 2012-07-01T09:55:31Z
.
. [skipping some output]
.
[etc.]
```

B   Appendix B: Python class for basic, space-naive Levenshtein implementation

B.1   Code

```python
import sys

class LevDistBasic:
    e = [] #edit operation array
    t = [] #grid array
    x = "" #string1
    y = "" #string2
    m = 0 #length string1
    n = 0 #length string2
    dist = 0 #Levenshtein distance
    ed = [] #the edit operation, calculated in _calculate()
    isFile = False

    def __init__(self, _x, _y, isFile=False):
        self.x = self._variablehandle(_x)
        self.y = self._variablehandle(_y)
        self.m = len(self.x)
        self.n = len(self.y)
        self.t = [[0]*(self.n+1) for _ in xrange(self.m+1)]
        self.e = [[" "]*(self.n+1) for _ in xrange(self.m+1)]
        self.dist = self._calculate()

    def __str__(self):
        return str(self.distance())

    def distance(self):
        return self.dist

    def strings(self):
        return self.x, self.y

    def table(self):
        return self.t

    def operation(self):
        return self.ed

    ##ADD WARNING for long strings / deal with them
    def showtable(self):
        result = ""
        for ch in self.y:
            result = result + ch + " "
        print "     ", result
        for r in range(len(self.t)):
            s = ' '
            if r:
                s = self.x[r-1]
            print s, ' ', self.t[r]

    def showop(self):
        for i, op in enumerate(self.ed):
            l = str(i) + ": "
            if op[0] == 'I':
                l += "insert " + op[-1]
```

11

```python
            elif op[0] == 'K':
                l += "keep " + op[-1]
            elif op[0] == 'D':
                l += "delete " + op[-1]
            elif op[0] == 'S':
                l += "swap " + op[-1][0] + " for " + op[-1][-1]
            else:
                return "FAIL: incorrect operation"
            print l

    def _ed(self):
        i, j = len(self.e)-1, len(self.e[0])-1
        self._ed_recursive(i,j)

    def _ed_recursive(self,i,j):
        if self.e[i][j] == ' ':
            if i == 0 and j > 0:
                self.ed.append(('D', self.y[0]))
            if j == 0 and i > 0:
                self.ed.append(('D', self.x[0]))
            return
        if self.e[i][j] == 'K':
            self._ed_recursive(i-1, j-1)
            self.ed.append((self.e[i][j], self.x[i-1]))
        elif self.e[i][j] == 'S':
            self._ed_recursive(i-1, j-1)
            self.ed.append((self.e[i][j], (self.x[i-1] + ',' + self.y[j-1])))
        elif self.e[i][j] == 'D':
            self._ed_recursive(i-1,j)
            self.ed.append((self.e[i][j], self.x[i-1]))
        else:
            self._ed_recursive(i,j-1)
            self.ed.append((self.e[i][j], self.y[j-1]))

    def _variablehandle(self,v):
        if not isinstance(v, str):
            try:
                return v.read()
            except:
                try:
                    return str(v)
                except:
                    print "Argument cannot be of type" + type(v)
                    raise
                pass
        return v

    def _calculate(self):
        for i in xrange(self.m+1):
            self.t[i][0] = i
        for j in xrange(self.n+1):
            self.t[0][j] = j
        j = 1
        while j < self.n+1:
            i = 1
            while i < self.m+1:
                c = (self.x[i-1] != self.y[j-1])
                dl = self.t[i-1][j] + 1
                ins = self.t[i][j-1] + 1
                sbs = self.t[i-1][j-1] + c
                self.t[i][j] = min(ins, dl, sbs)
```

```
                if ins < dl and ins < sbs:
                    self.e[i][j] = 'I'
                elif dl <= sbs:
                    self.e[i][j] = 'D'
                else:
                    if(self.x[i-1] != self.y[j-1]):
                        self.e[i][j] = 'S'
                    else:
                        self.e[i][j] = 'K'
                i += 1
            j += 1
        self._ed()
        return self.t[self.m][self.n]
```

B.2   Example output

```
$ python
>>> import LevDistBasic
>>> test = LevDistBasic.LevDistBasic("bank","book")
>>> test.showtable()


        b   o   o   k
    [0, 1, 2, 3, 4]
b   [1, 0, 1, 2, 3]
a   [2, 1, 1, 2, 3]
n   [3, 2, 2, 2, 3]
k   [4, 3, 3, 3, 2]


>>> t = test.table()
>>> print t

[[0, 1, 2, 3, 4], [1, 0, 1, 2, 3], [2, 1, 1, 2, 3], [3, 2, 2, 2, 3], [4, 3, 3, 3, 2]]

>>> s = test.strings()
>>> print s

('bank', 'book')

>>> test.showop()

0: keep b
1: swap a for o
2: swap n for o
3: keep k

>>> ed = test.operation()
>>> print ed

[('K', 'b'), ('S', 'a,o'), ('S', 'n,o'), ('K', 'k')]


>>> print test

2
```

References

On wikipedia

AcaWiki. http://acawiki.org. Accessed: 2014-04-31 (see p. 2).

Adler, B. Thomas and Luca de Alfaro (2007). "A Content-driven Reputation System for the Wikipedia". In: *Proceedings of the 16th International Conference on World Wide Web*. WWW '07. Banff, Alberta, Canada: ACM, pp. 261–270. ISBN: 978-1-59593-654-7. DOI: 10.1145/1242572.1242608. URL: http://doi.acm.org/10.1145/1242572.1242608 (see p. 4).

Alexa About Us. http://www.alexa.com/about. Accessed: 2014-04-31 (see p. 2).

Alexa The top 500 sites on the web. http://www.alexa.com/topsites. Accessed: 2014-04-31 (updated daily) (see p. 2).

Cross, Tom (2006). "Puppy smoothies: Improving the reliability of open, collaborative wikis". In: *First Monday* 11.9. ISSN: 13960466. URL: http://firstmonday.org/ojs/index.php/fm/article/view/1400 (see p. 4).

De Alfaro, Luca, Ashutosh Kulshreshtha, Ian Pye, and B. Thomas Adler (Aug. 2011). "Reputation Systems for Open Collaboration". In: *Commun. ACM* 54.8, pp. 81–87. ISSN: 0001-0782. DOI: 10.1145/1978542.1978560. URL: http://doi.acm.org/10.1145/1978542.1978560 (see p. 4).

Denning, Peter, Jim Horning, David Parnas, and Lauren Weinstein (Dec. 2005). "Wikipedia Risks". In: *Commun. ACM* 48.12, pp. 152–152. ISSN: 0001-0782. DOI: 10.1145/1101779.1101804. URL: http://doi.acm.org/10.1145/1101779.1101804 (see pp. 2, 3).

Giles, Jim (2005). "Internet encyclopaedias go head to head". In: *Nature* 438.7070, pp. 900–901. ISSN: 0028-0836. URL: http://dx.doi.org/10.1038/438900a (see pp. 3, 4).

Kittur, Aniket, Bongwon Suh, Bryan A. Pendleton, and Ed H. Chi (2007). "He Says, She Says: Conflict and Coordination in Wikipedia". In: *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*. CHI '07. San Jose, California, USA: ACM, pp. 453–462. ISBN: 978-1-59593-593-9. DOI: 10.1145/1240624.1240698. URL: http://doi.acm.org/10.1145/1240624.1240698 (see pp. 2–4).

Lieberman, Michael and Jimmy Lin (2009). *You Are Where You Edit: Locating Wikipedia Contributors through Edit Histories*. URL: http://aaai.org/ocs/index.php/ICWSM/09/paper/view/204 (see p. 5).

Lih, Andrew (2004). "Wikipedia as Participatory journalism: reliable sources? metrics for evaluating collaborative media as a news resource". In: *In Proceedings of the 5th International Symposium on Online Journalism*, pp. 16–17 (see pp. 3, 4).

Lucassen, Teun and Jan Schraagen (2011). "Evaluating WikiTrust: A trust support tool for Wikipedia". In: *First Monday* 16.5. ISSN: 13960466. URL: http://firstmonday.org/ojs/index.php/fm/article/view/3070 (see p. 4).

Mesgari, Mostafa, Chitu Okoli, Mohamad Mehdi, Finn Årup Nielsen, and Arto Lanamäki (2014). ""The sum of all human knowledge": A systematic review of scholarly research on the content of Wikipedia". In: *Journal of the American Society for Information Science and Technology*. This is a postprint of an article accepted for publication in Journal of the American Society for Information Science and Technology copyright © 2014 (American Society for Information Science and Technology). URL: http://spectrum.library.concordia.ca/978618/ (see pp. 2, 15).

Mestyån, Mårton, Taha Yasseri, and Jånos Kertåsz (2012). "Early Prediction of Movie Box Office Success based on Wikipedia Activity Big Data". In: *CoRR* abs/1211.0970. URL: `http://dblp.uni-trier.de/db/journals/corr/corr1211.html#abs-1211-0970` (see p. 5).

Potthast, Martin, Benno Stein, and Robert Gerling (2008). "Automatic Vandalism Detection in Wikipedia". In: *Advances in Information Retrieval*. Ed. by Craig Macdonald, Iadh Ounis, Vassilis Plachouras, Ian Ruthven, and RyenW. White. Vol. 4956. Lecture Notes in Computer Science. Springer Berlin Heidelberg, pp. 663–668. ISBN: 978-3-540-78645-0. DOI: `10.1007/978-3-540-78646-7_75`. URL: `http://dx.doi.org/10.1007/978-3-540-78646-7_75` (see p. 3).

Python wikipedia 1.3.0. `https://pypi.python.org/pypi/wikipedia`. "Wikipedia is a Python library that makes it easy to access and parse data from Wikipedia." Accessed: 2014-04-31 (see p. 7).

Stvilia, B., M.B. Twidale, L.C. Smith, and L. Gasser (2005). "Assessing information quality of a community-based encyclopedia". In: *Proceedings of the International Conference on Information Quality*, pp. 442–454 (see p. 4).

Suh, Bongwon, Gregorio Convertino, Ed H. Chi, and Peter Pirolli (2009). "The Singularity is Not Near: Slowing Growth of Wikipedia". In: *Proceedings of the 5th International Symposium on Wikis and Open Collaboration*. WikiSym '09. Orlando, Florida: ACM, 8:1–8:10. ISBN: 978-1-60558-730-1. DOI: `10.1145/1641309.1641322`. URL: `http://doi.acm.org/10.1145/1641309.1641322` (see p. 2).

WikiLit. `http://wikilit.referata.com`. A temporary site, relating directly to[42]. Accessed: 2014-04-31 (see p. 2).

Wikimedia enwiki dump progress on 20140502. `http://dumps.wikimedia.org/enwiki/20140502/`. Accessed: 2014-04-29 (see p. 3).

Wikimedia Wikipedia.org is more popular than...: Note on Alexa rankings. `https://meta.wikimedia.org/wiki/Wikipedia.org_is_more_popular_than...#Note_on_Alexa_rankings`. Accessed: 2014-04-31 (see p. 2).

WikiPapers. `http://wikipapers.referata.com`. Accessed: 2014-04-31 (see p. 2).

Wikipedia Counter-Vandalism Unit. `https://en.wikipedia.org/wiki/Wikipedia:CVU`. Accessed: 2014-04-31 (see p. 3).

Wikipedia Wikipedia Signpost/2013-02-04/Special report: Examining the popularity of Wikipedia articles: catalysts, trends, and applications. `https://en.wikipedia.org/wiki/Wikipedia:Wikipedia_Signpost/2013-02-04/Special_report`. Accessed: 2014-04-31 (see p. 3).

Wilkinson, Dennis M. and Bernardo A. Huberman (2007). "Assessing the Value of Coooperation in Wikipedia". In: *CoRR* abs/cs/0702140 (see p. 2).

Zeng, Honglei, Maher A. Alhossaini, Li Ding, Richard Fikes, and Deborah L. McGuinness (2006). "Computing Trust from Revision History". In: *Proceedings of the 2006 International Conference on Privacy, Security and Trust: Bridge the Gap Between PST Technologies and Business Services*. PST '06. Markham, Ontario, Canada: ACM, 8:1–8:1. ISBN: 1-59593-604-1. DOI: `10.1145/1501434.1501445`. URL: `http://doi.acm.org/10.1145/1501434.1501445` (see p. 4).

---

[42]Mesgari, Okoli, Mehdi, Nielsen, and Lanamäki 2014.

On edit distance in general

Levenshtein, VI (1966). "Binary Codes Capable of Correcting Deletions, Insertions and Reversals". In: *Soviet Physics Doklady* 10, p. 707 (see p. 4).